

Thread (II)




User Threads & User/Kernel Mapping

User Threads





- Created by a thread library and scheduling is managed by the library itself in user space (the existence of user threads is unknown to the OS)

User Thread Libraries

Pthreads (POSIX Threads)

-  A POSIX standard (IEEE 1003.1c) API for thread creation and synchronization
-  API specifies behavior of the thread library, implementation is up to development of the library
-  Common in UNIX operating systems (Solaris, Linux, Mac OS X)

Java Threads

-  Managed by the JVM
-  May be created by:
 -  Extending Thread class
 -  Implementing the Runnable interface

Win32 Threads

-  Similar to pthread, differing in function names

PThreads

```
# include <pthread.h>
```

```
✧ int pthread_create ( pthread_t * thread,  
                      pthread_attr_t * attr,  
                      void * (*start_routine)(void *),  
                      void *arg )
```

- ✧ Creates a new thread
- ✧ 1st argument: the address of the created thread
- ✧ 2nd argument: the address of the structure containing the attributes of the created thread
 - ✧ For default attribute values, make this argument as ***NULL***
- ✧ 3rd argument: the pointer to the function that is the code for the created thread
- ✧ 4th argument: the pointer to the argument of the start_routine function. If multiple arguments are needed, define a data structure that contains all arguments.

PThreads

include <pthread.h>

✧ void **pthread_exit** (void *status)

✧ Terminate execution of the calling thread

✧ 1st argument: Points to an optional termination status. If no termination status is desired, its value should be NULL.

✧ pthread_t **pthread_self** (void)


✧ returns the *pthread* handle (ID) of the calling thread

PThreads

include <pthread.h>

- ✧ int **pthread_join** (pthread_t tid, void **status)
 - ✧ "Joining" is one way to accomplish synchronization between threads.
 - ✧ Blocks the caller until the specified thread terminates
 - ✧ 1st argument: the *id* of the thread to be waited
 - ✧ 2nd argument: the address of the variable to receive the thread's exit status
 - ✧ Usually set to ***NULL*** or ***int****

Example: Hello World!

 This simple example code creates 5 threads with the `pthread_create()` routine. Each thread prints a "Hello World!" message, and then terminates with a call to `pthread_exit()`.

```
#include <pthread.h>
#include <stdio.h>
#define NUM_THREADS    5

void *PrintHello(void *threadid)
{
    printf("\n%d: Hello World from Thread #%d !\n",
          threadid, pthread_self());
    pthread_exit(NULL);
}
```

Example: Hello World! (con't)

```
int main (int argc, char *argv[])
{
    pthread_t threads[NUM_THREADS];
    int rc, t;
    for (t=0;t < NUM_THREADS;t++) {
        printf("Creating thread %d\n", t);
        rc = pthread_create(&threads[t], NULL, PrintHello, (void *) t);
        if (rc){
            printf("ERROR; return code from pthread_create() is %d\n", rc);
            exit(-1);
        }
    }

    for(t=0; t<NUM_THREADS;t++) pthread_join(threads[t],NULL);
    pthread_exit(NULL);
}
```


Example: Compilation

Suppose the program file is named: HelloWorld.c

```
$gcc -o HelloWorld HelloWorld.c -lpthread
```

Java Threads

- Defining a class X that implements the **Runnable** interface to contain the implementation of a thread

```
public interface Runnable
{
    public abstract void run();
}
```

- Creating a Thread object to wrap the class X
- Manipulate the Thread object by calling its methods
 - run() method
 - join() method
 -

Example: Summation

```
class Summation implements Runnable
```

```
{
```

```
    private int upper;
```

```
    public int sum;
```

```
    public Summation(int upper, int sum) {
```

```
        this.upper = upper;
```

```
        this.sum = sum;
```

```
    }
```

```
    public void run() {
```

```
        int sum = 0;
```

```
        for (int i=0; i<=upper; i++) sum += i;
```

```
        System.out.println("Sum of 1 through " + upper + " is " + sum);
```

```
    }
```

```
}
```

Example: Summation

```
public class Driver{
    public static void main(String[ ] args) {
        if (args.length >0) {
            if (Integer.parseInt(args[0])<0)
                System.err.println(args[0]+"must be >=0");
            else{
                int upper = Integer.parseInt(args[0]);
                Summation sumObj = new Summation(upper,0);
                Thread thrd=new Thread(sumObj);
                thrd.start( );
                try {thrd.join( );} catch (InterruptedException ie) { }
            }
        }
    }
}
```

Example: Summation

```
$javac Summation.java
```

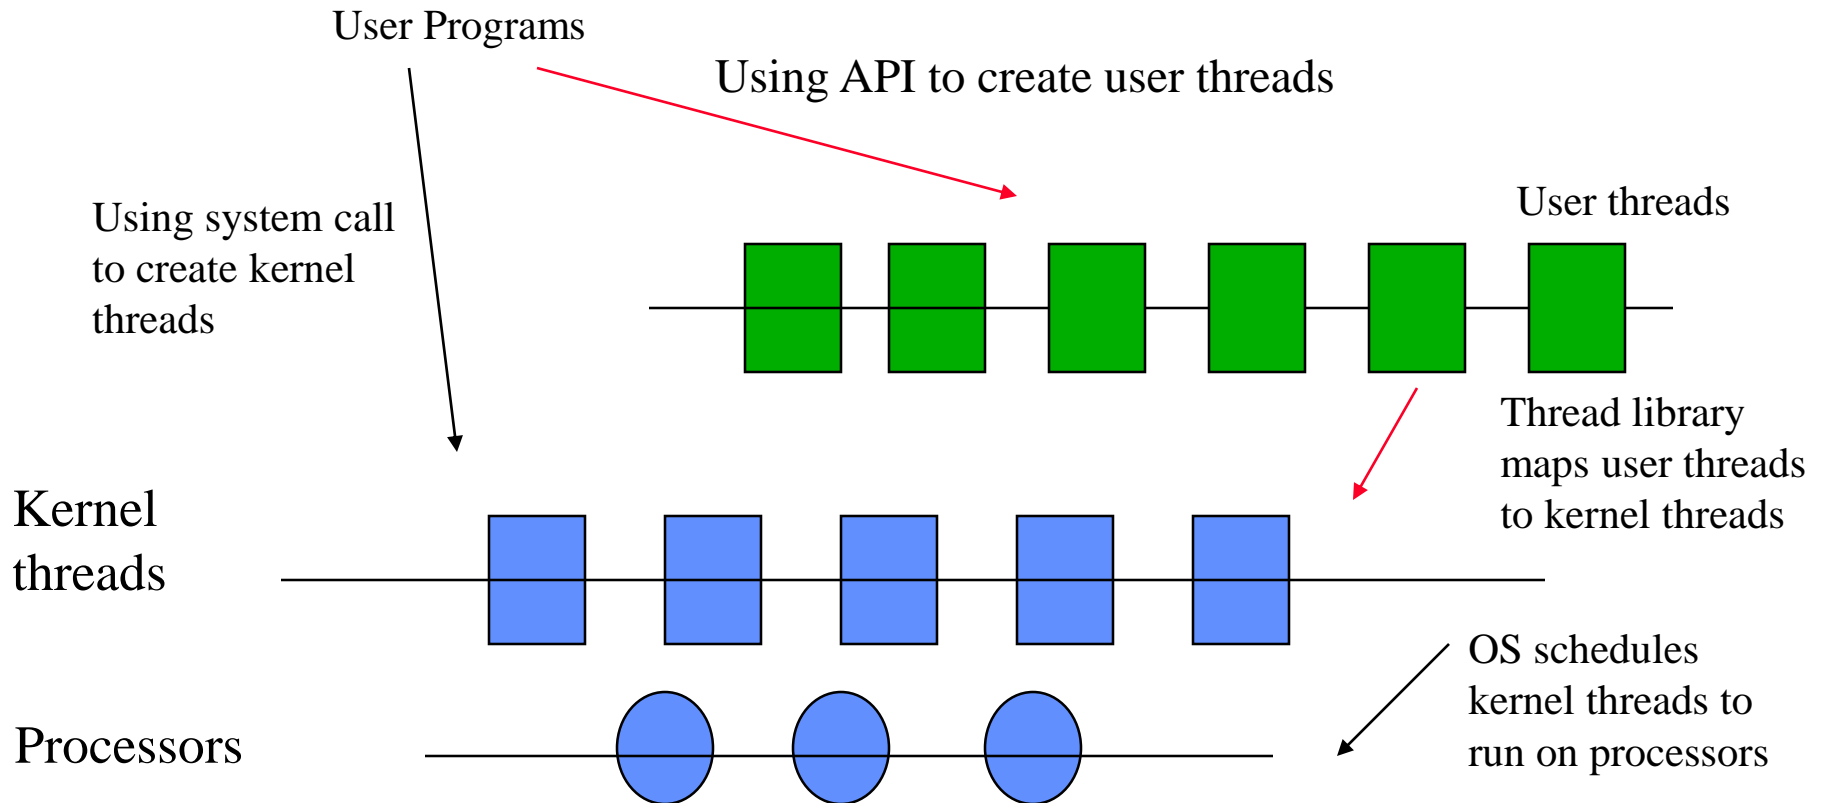
```
$javac Driver.java
```

```
$java Driver 100
```

```
Sum of 1 through 100 is 5050
```

Threads in a Computer System

- A thread library provides the programmer with an API for creating and managing threads.



Kernel Threads vs. User Threads

- Kernel threads:

- Directly created/managed by the OS kernel

- User threads




- created by a thread library and scheduling is managed by the library itself (the existence of user threads is unknown to the OS)

- low-cost in thread creation; portable


- may not utilize multi-processors efficiently (depends on mapping strategy between user and kernel threads)

- When executed, must be mapped to kernel threads

Mapping User Threads to Kernel Threads

-  Many to One
-  One to One
-  Many to Many

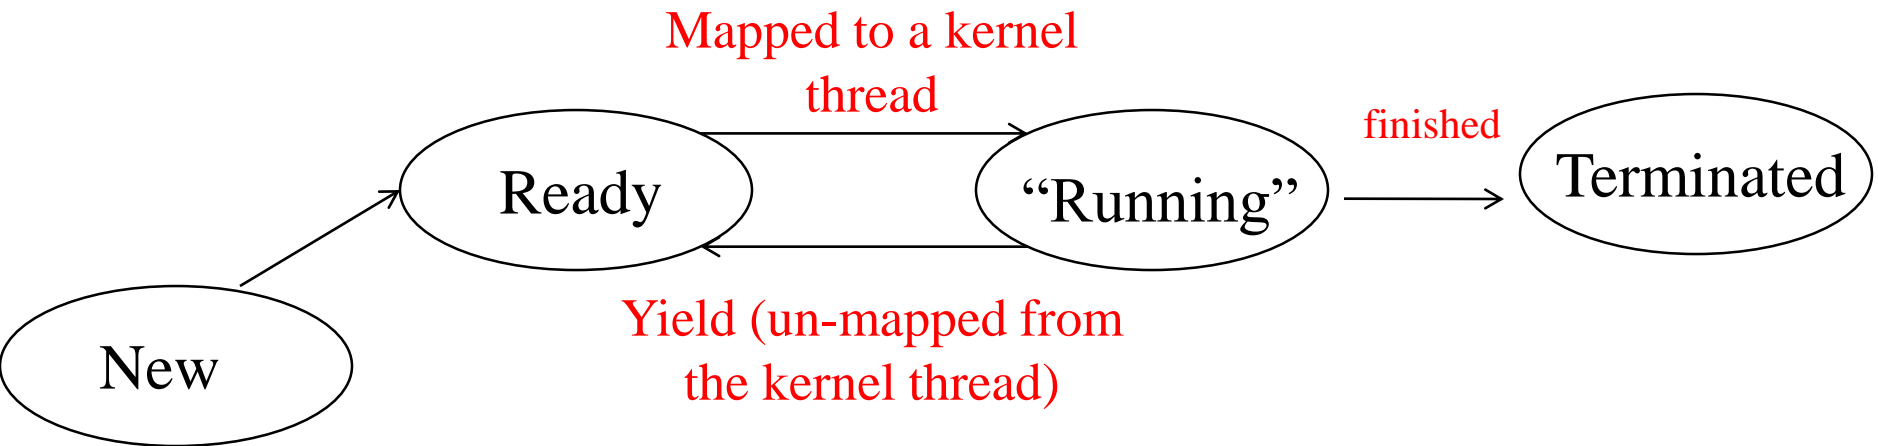
Many-to-One



- Many user-level threads within the same process are mapped to single kernel thread
- User thread library schedules user threads to the kernel thread
- The multiple threads “time-share” the single kernel thread
 - A kernel thread is created for a process
 - Upon the thread yields, switch to another thread in the same process (kernel is unaware of it)
 - Upon the thread blocks, all threads in the process block

Thread Scheduling in User Space

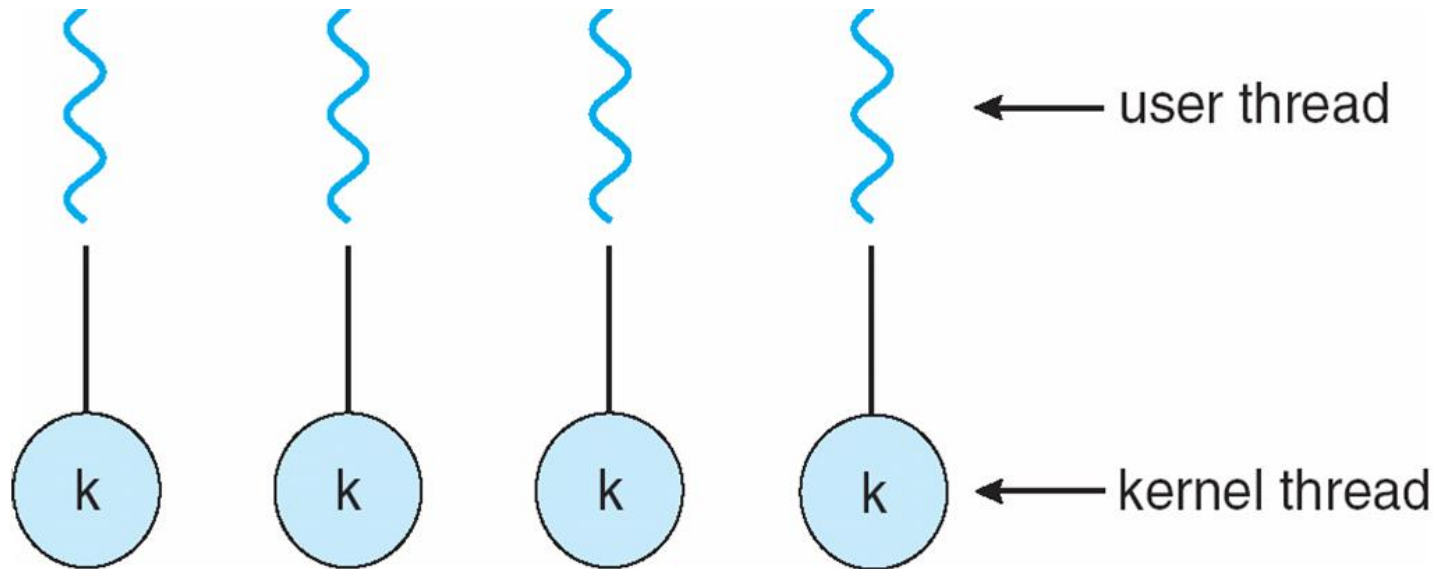
- 📖 User-space scheduling: managing the (dynamic) mapping from user threads to kernel threads.



Lifetime of a user thread

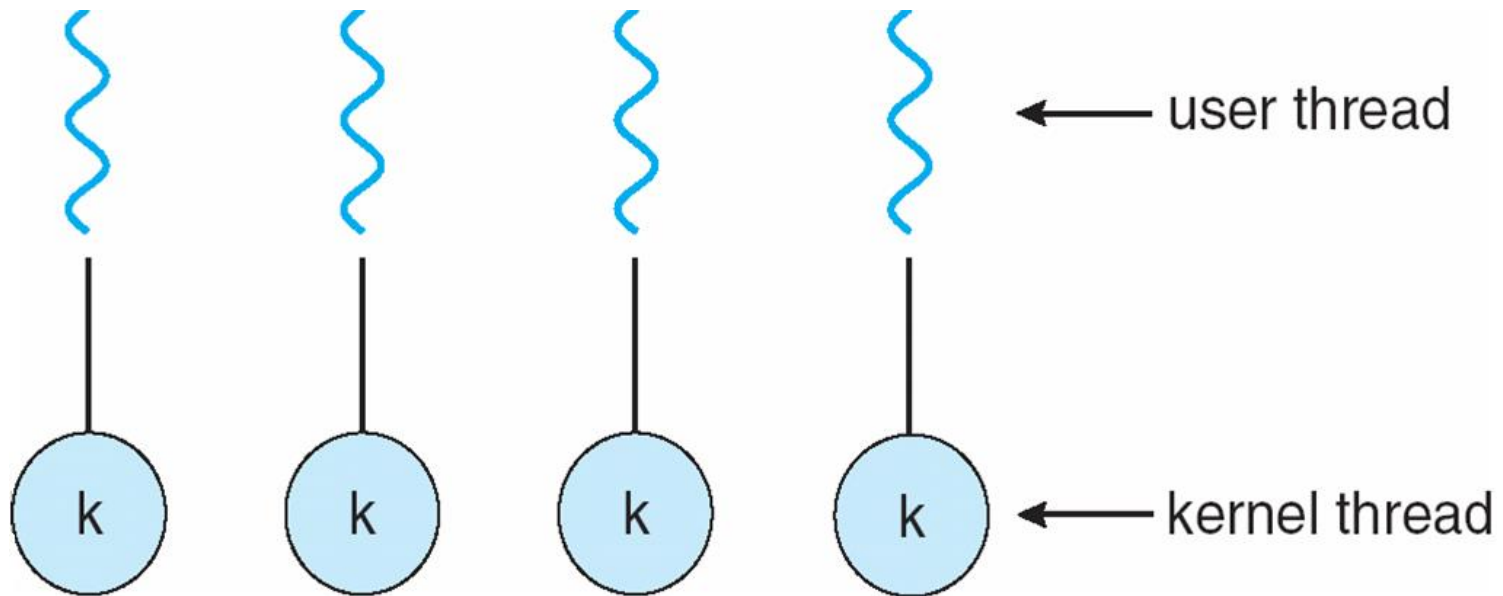
One-to-One

- Each user-level thread maps to a kernel thread
- User thread library simply provides a portable interface for thread creation/management, of which the implementation rely on the kernel thread (kernel of the OS)

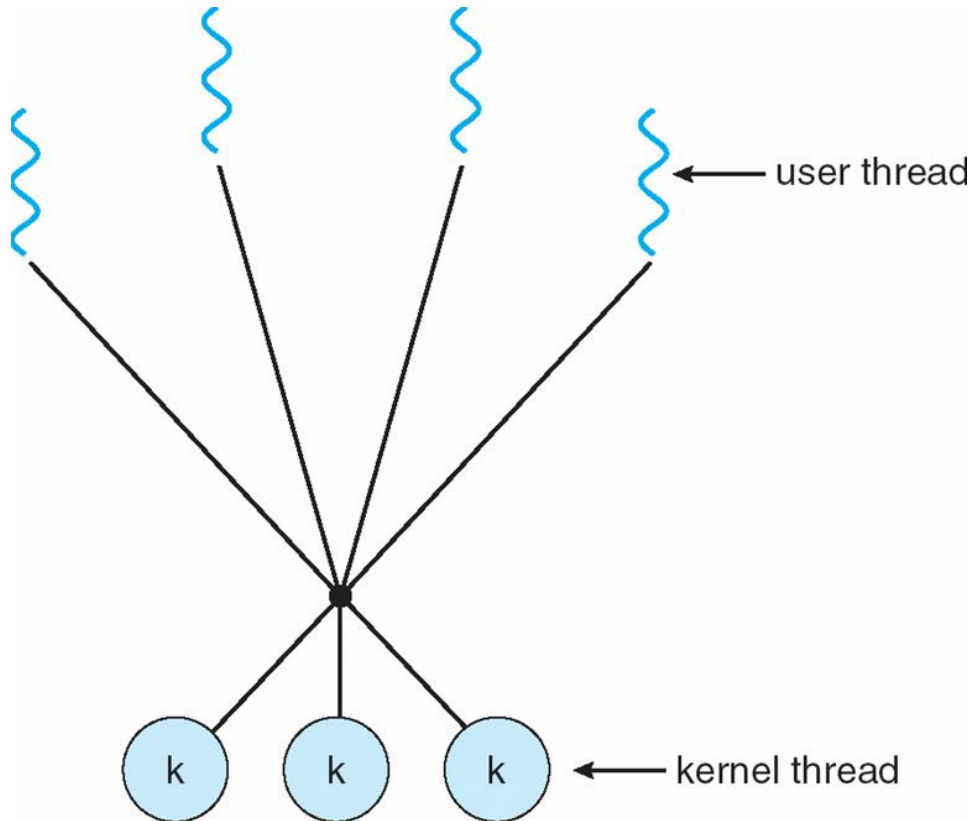


One-to-One

- Upon a thread currently running on a CPU yields or blocks, CPU can be switched to another thread (kernel is aware of this)

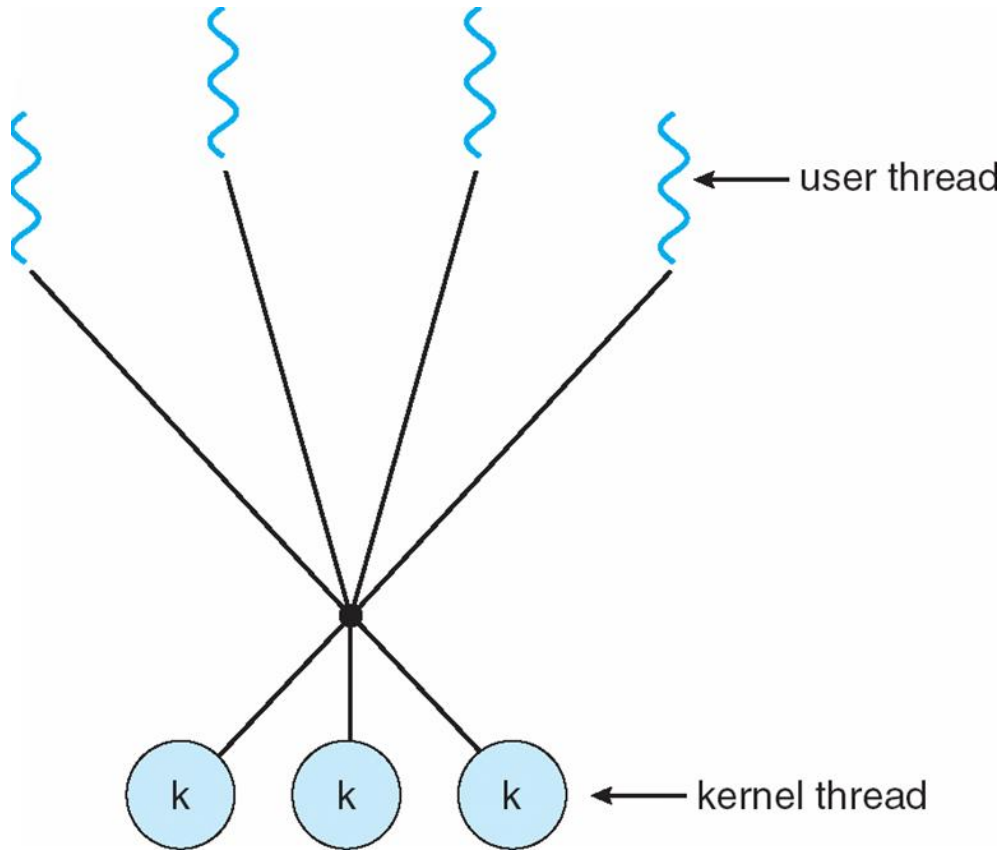


Many-to-Many Model




- Allows many user level threads to be mapped to many kernel threads
- User thread library schedules user threads to kernel threads

Many-to-Many Model



- Upon a currently-running user thread yields, another thread may be mapped to the yielding user thread's kernel thread and runs
- Upon a currently-running user thread blocks, other thread may run on the CPU


Two-level Model





- 📖 A sub-type of M:M
- 📖 It allows a user thread (e.g., with high priority) to be **bound** to a kernel thread

Thread Context and User/Kernel Thread Mapping

Thread context

-  Each thread has a “context”: User Stack + Thread Control Block (TCB: register values, etc.)

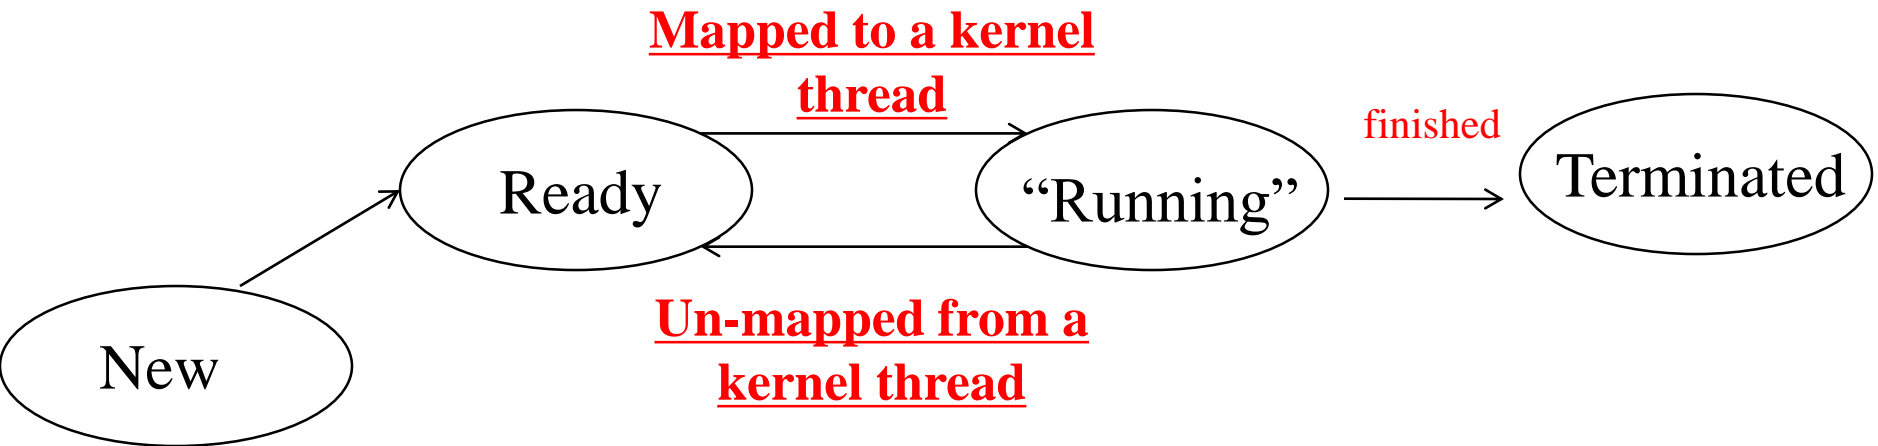
-  A kernel thread's context is managed by the OS kernel

-  A user thread's context is managed by the user thread library

-  Mapping a user thread to a kernel thread is achieved by copying the user thread's context to the kernel thread's context.

Thread Scheduling in User Space

- ❏ User-space scheduling: managing the (dynamic) mapping from user threads to kernel threads.



Lifetime of a user thread

Linux Supports for Thread Context Management

- ❏ Data structure `ucontext_t`

- ❏ System calls

 - ❏ `getcontext`

 - ❏ `makecontext`

 - ❏ `setcontext`

 - ❏ `swapcontext`

- ❏ All declared in `ucontext.h`

Data Type “ucontext_t”

```
typedef struct ucontext{  
    struct ucontext    *uc_link;  
        //points to the context that will be restored  
        //when the current one terminates  
    ...  
    stack_t            uc_stack;  
        //the stack used by this context  
    ... register value, etc. ...  
}ucontext_t;
```

System Call “getcontext”

```
int getcontext(ucontext_t *ucp)
```

To [initialize](#) the structure pointed by ucp to contain a copy of the currently-active kernel thread's context

```
#include <ucontext.h>
```

```
...
```

```
ucontext_t *ucp;
```

```
ucp = (ucontext_t *)malloc(sizeof(ucontext_t));
```

```
...
```

```
getcontext(ucp);
```

```
...
```

System Call “makecontext”

`void makecontext(ucontext_t *ucp, void (*func)())`

To make a thread context (stored in the structure pointed by ucp) for function func(); the resulting context becomes the context for a thread running this func().

```
#include <ucontext.h>
```

```
void func() {...}
```

```
ucontext_t *ucp;
```

```
ucp = (ucontext_t *)malloc(sizeof(ucontext_t));
```

```
getcontext(ucp); //initialize the structure
```





```
ucp->uc_stack.ss_sp = malloc(16384); //create a new stack space for  
the new thread
```

```
ucp->uc_stack.ss_size = 16384;
```

```
makecontext(ucp, func);
```

System Call “setcontext”

```
int setcontext(ucontext_t *ucp)
```

-  The function updates the context of the currently-running thread with the context pointed by ucp; as the effect, the thread with context *ucp starts/resumes its execution.
-  This can be used to implement mapping from a user thread (thread with context *ucp) to a kernel thread (currently active one)!
-  A successful call does not return.
-  The context should be a valid one (i.e., obtained by a call of context or make context)

System Call “setcontext”

```
...
int cloned_func() {
    ...
    setcontext(ucp); //user thread with context ucp is run
                      //(i.e., the user thread is mapped to this kernel thread
    ...
}
...
main() {
    ...
    clone(cloned_func, ...); //creates a kernel thread
    ...
}
```

System Call “swapcontext”

```
int swapcontext(ucontext_t *oucp, ucontext_t *ucp)
```

- ❏ This function saves the context of the currently-active kernel thread to the structure pointed by oucp (you should have already allocated the space for the structure); as the effect, the thread with context *oucp is started or resumed.