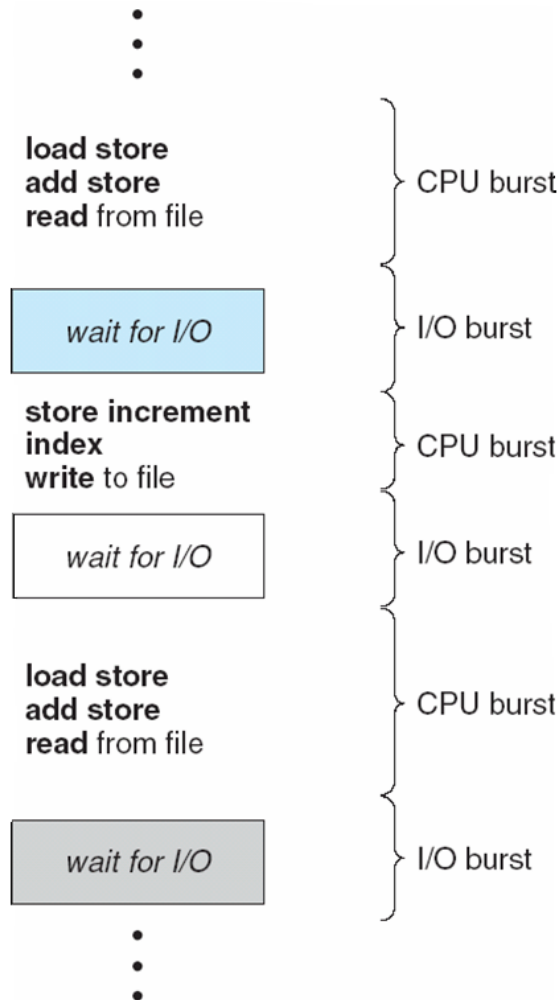


CPU Scheduling (I)

CPU Scheduling



☞ Decide which process should be run on the CPU






☞ Purpose: To maximize CPU utilization obtained with multiprogramming/multitasking

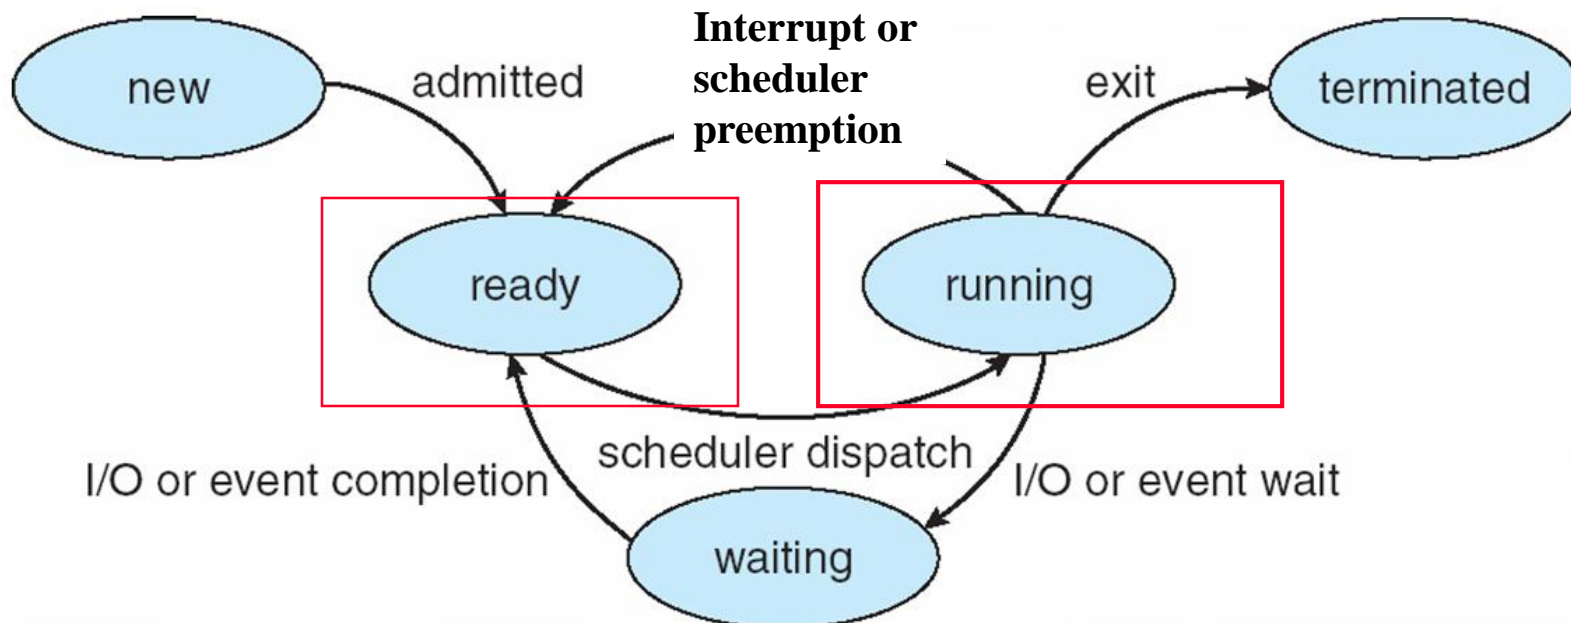
☞ Model of a process: CPU–I/O Burst Cycle – Process execution consists of a *cycle* of CPU execution and I/O wait


Notes

- ❏ In this discussion of CPU scheduling, we use the term “process” to represent a “traditional” process of only a single thread.
- ❏ The principles and algorithms also apply to kernel threads, if the OS supports kernel threads. Note: if a process is allowed to have multiple kernel threads, it is the kernel threads (not processes) that are scheduled by the OS.

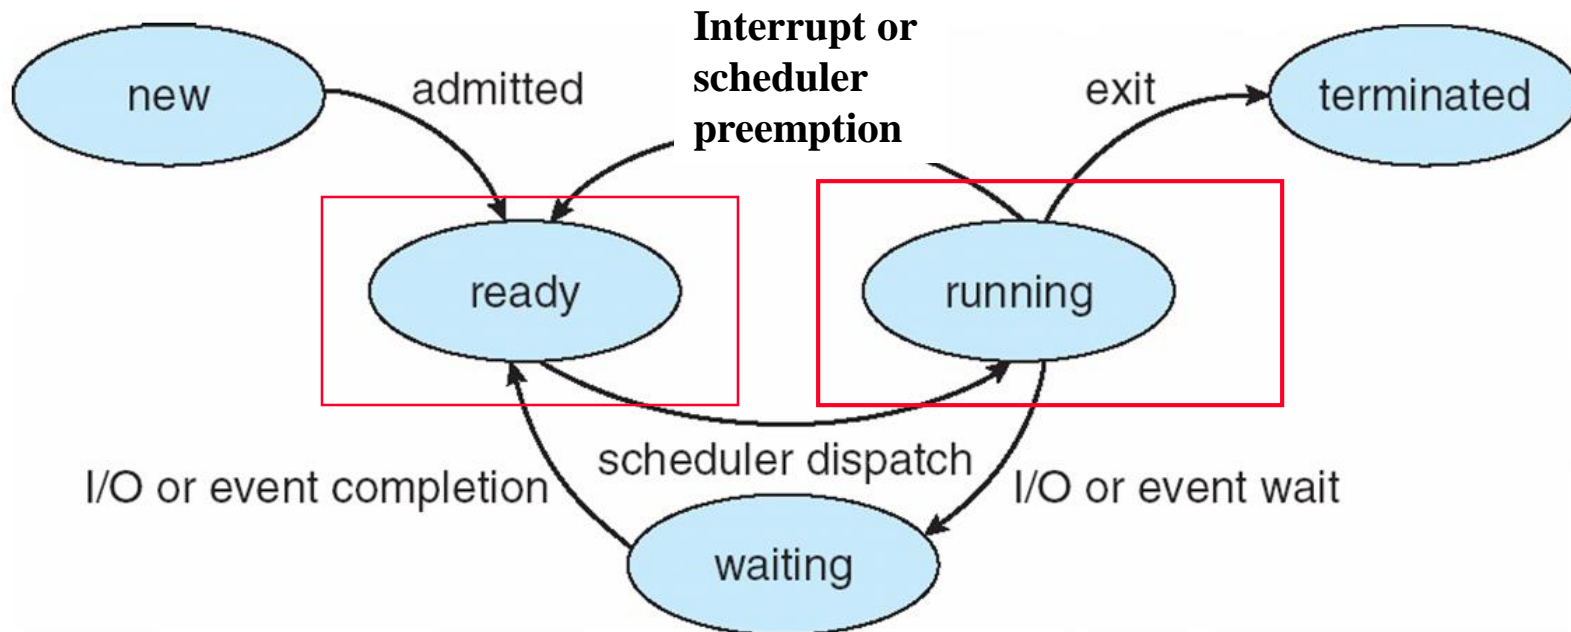
Recall: Process State and Transitions

-  **new:** The process is being created
-  **running:** Instructions are being executed
-  **waiting:** The process is waiting for some event to occur
-  **ready:** The process is waiting to be assigned to a processor
-  **terminated:** The process has finished execution





 CPU scheduling decisions may take place when a process:


1. Switches from running to waiting state
2. Switches from running to ready state
3. Switches from **waiting/new** to ready
4. Terminates




CPU Scheduler






-  Selects from among the processes in memory that are ready to execute, and allocates the CPU to one of them

-  CPU scheduling decisions may take place when a process:
 1. Switches from running to waiting state
 2. Switches from running to ready state
 3. Switches from **waiting/new** to ready
 4. Terminates

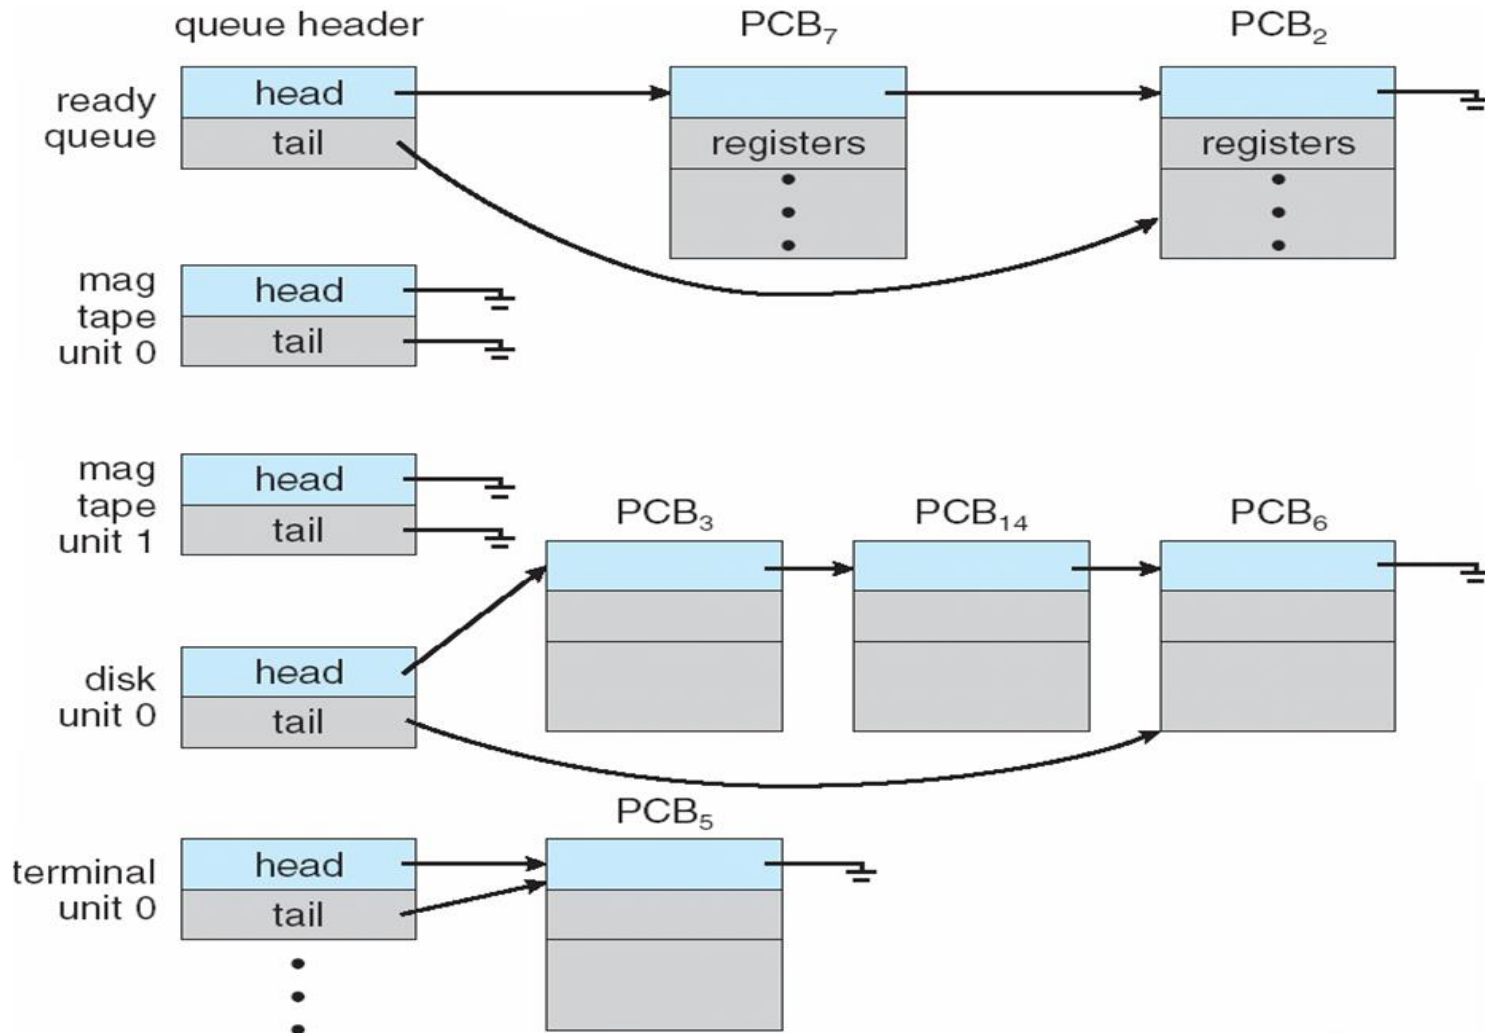
-  Scheduling under 1 and 4 is **nonpreemptive**

-  All other scheduling is **preemptive**


Process Queues

-  **Job queue** – the queue of the PCBs of all the processes in the system
-  **Ready queue** – the queue of the PCBs of all processes residing in main memory, ready and waiting to execute
-  **Device queues** – the queue of PCBs of processes waiting for an I/O device
-  Processes (i.e., their PCBs) migrate among the various queues
-  **Note:** When it is threads that are scheduled, the queues are composed of TCBs.

Ready Queue And Various I/O Device Queues



Context Switch

 When CPU switches from one process to another one, the system (hardware/OS) must save the state of the old process and load the saved state of the new one via a **context switch**

Context of a Running Process

- ☐ User memory space

- ☐ Value of “program counter register”

 - ☐ storing the address of the instruction to be executed next

- ☐ Value of “stack pointer register”

 - ☐ pointing to the top of current stack

- ☐ Values of other registers

 - ☐ storing intermediate results of ongoing computations

(When a process is running, the above information is stored on the registers.)

- ☐ Process state, memory info, I/O info, ... : stored in PCB

Context of a Non-Running Process

- ☐ User memory space

- ☐ Value of “program counter register”

 - ☐ storing the address of the instruction to be executed next

- ☐ Value of “stack pointer register”

 - ☐ pointing to the top of current stack

- ☐ Values of other registers

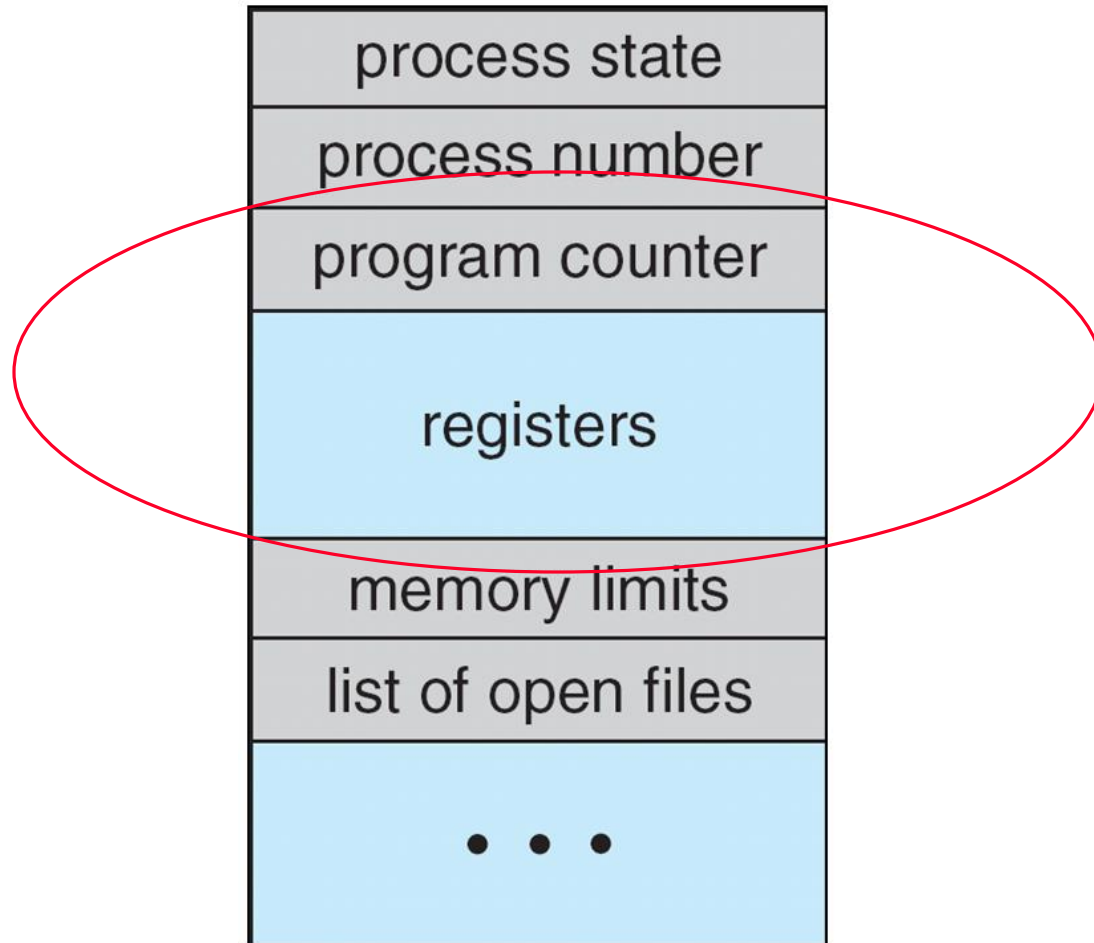
 - ☐ storing intermediate results of ongoing computations

(When a process is not running, the above information is stored in PCB.)

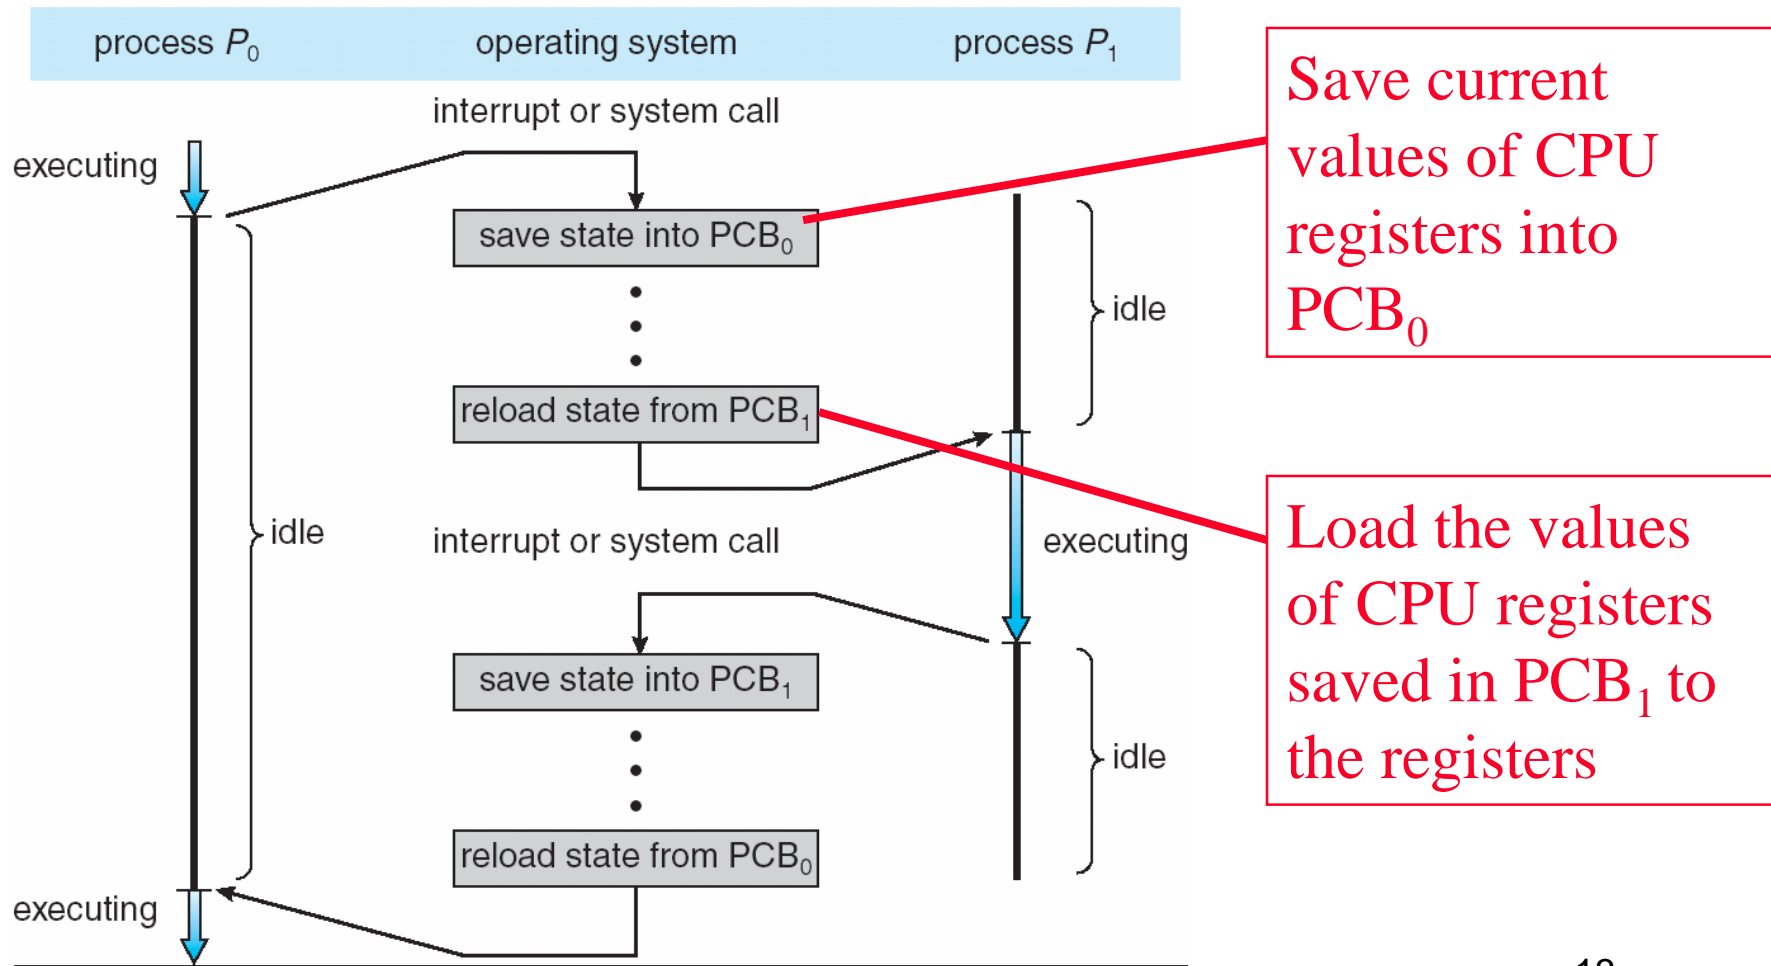
- ☐ Process state, memory info, I/O info, ... : stored in PCB

Note: for a thread, it is TCB. 11

Process Control Block (PCB)



CPU Switch between Processes








Context Switch

- ❏ Context-switch time is overhead
- ❏ Time depends on hardware support
 - ❏ memory speed;
 - ❏ number of registers whose values should be copied;
 - ❏ the existence of special instructions for storing/loading all registers
- ❏ Typically, it takes a few microseconds

Evaluation Metrics

- ❏ **CPU utilization** – the fraction of valued CPU time over the whole time window
- ❏ **Throughput** – # of processes that complete their execution per time unit
- ❏ **Turnaround time** – amount of time to execute a particular process (from when a process is created/submitted to when it completes)
- ❏ **Waiting time** – amount of time a process has been waiting in the ready queue
- ❏ **Response time** – amount of time it takes from when a request was submitted until the first response is produced (for time-sharing environment)

Scheduling Algorithm Optimization Criteria

-  Maximize CPU utilization
-  Maximize throughput
-  Minimize turnaround time
-  Minimize waiting time
-  Minimize response time

First-Come, First-Served (FCFS) Scheduling


<u>(Ready) Process</u>	<u>(Next) Burst Time</u>
------------------------	--------------------------

P_1	24
-------	----

P_2	3
-------	---

P_3	3
-------	---

❏ Suppose that the processes arrive in the order: P_1, P_2, P_3
The time interval between arrivals is negligible. The **Gantt Chart** for the schedule is:



❏ Waiting time for $P_1 = 0$; $P_2 = 24$; $P_3 = 27$

❏ Average waiting time: $(0 + 24 + 27)/3 = 17$

FCFS Scheduling (Cont)

Suppose that the processes arrive in the order

$$P_2, P_3, P_1$$

📖 The Gantt chart for the schedule is:



📖 Waiting time for $P_1 = 6$, $P_2 = 0$, $P_3 = 3$

📖 Average waiting time: $(6 + 0 + 3)/3 = 3$

📖 Much better than previous case


📖 *Convoy effect* when short processes behind a long process


Shortest-Job-First (SJF) Scheduling

 Schedule first the process with the shortest burst time

<u>Process</u>	(CPU) <u>Burst Time</u>
P_1	6
P_2	8
P_3	7
P_4	3

 SJF scheduling chart



 Average waiting time = $(3 + 16 + 9 + 0) / 4 = 7$

SJF: Preemptive or Nonpreemptive


- ❏ SJF can be either preemptive or nonpreemptive
 - ❏ Preemptive (Shortest-remaining-time-first): the currently running process can be preempted by a newly arriving process that has shorter CPU burst
 - ❏ Non-preemptive: even a newly arriving process has shorter CPU burst than the currently running one, the running process is not preemptive.

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

What are the Gantt charts for the preemptive SJF and the nonpreemptive SJF?²⁰


Non-preemptive SJF

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5



Preemptive SJF

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5




Priority Scheduling

- ❏ A priority number (integer) is associated with each process
- ❏ The CPU is allocated to the process with the highest priority (smallest integer \equiv highest priority)
 - ❏ Preemptive
 - ❏ Non-preemptive
- ❏ SJF is a priority scheduling where priority is the predicted next CPU burst time

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>	<u>Priority</u>
P_1	0	10	3
P_2	1	3	1
P_3	2	2	4
P_4	3	2	2


Priority Scheduling: Non-preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>	<u>Priority</u>
P_1	0	10	3
P_2	1	3	1
P_3	2	2	4
P_4	3	2	2



Priority Scheduling: Preemptive

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>	<u>Priority</u>
P_1	0	10	3
P_2	1	3	1
P_3	2	2	4
P_4	3	2	2



Priority Scheduling

- ❏ Problem: **Starvation** – low priority processes may never execute
- ❏ Solution: **Aging** – as time progresses increase the priority of the waiting process

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>	<u>Priority</u>
P_1	0	10	1
P_2	1	5	10
P_3	2	2	2
P_4	3	2	3
P_5	4	2	3
P_6	5	2	3
P_7	6	1	3
P_8	7	1	2
...