


Process Synchronization (II)

Semaphore


- ❏ An abstract data type
 - ❏ Semaphore S – integer variable
 - ❏ Two standard operations to access S : `wait()` and `signal()`
 - ❏ Originally called `P()` and `V()`
- ❏ Can only be accessed via two (atomic) operations

```
wait (S) {  
    while (S <= 0); //blocked  
    S--;  
}  
signal (S) {  
    S++;  
}
```

Usage of Semaphores

 **Binary** semaphore – integer value S can range only between 0 and 1

 Also known as **mutex** (i.e., mutual exclusive) locks

 Provides mutual exclusion (the following is pseudo code)

Semaphore mutex; // initialized to 1

wait (mutex);


... Critical Section ...

signal (mutex);

... remainder section ...

 **Counting** semaphore – integer value can range over an unrestricted domain

Dining-Philosophers Problem



Shared data

Bowl of rice (data set)

Semaphore **chopstick** [5] initialized to 1

Dining-Philosophers Problem

 The structure of Philosopher i :

```
wait ( chopstick[i] );  
wait ( chopStick[ (i + 1) % 5] );  
... eat...  
signal ( chopstick[i] );  
signal ( chopstick[ (i + 1) % 5] );  
... think...
```


POSIX Semaphores


- ❏ Two types of semaphores
 - ❏ Unnamed semaphores
 - ❏ Volatile
 - ❏ Storage is allocated by user program
 - ❏ Named semaphores
 - ❏ Permanent (like named Pipe (FIFO))
 - ❏ Storage is allocated by the OS
- ❏ Defined in `<semaphore.h>`
- ❏ Provided by Linux (You can use it on pyrite)

Unnamed Semaphores

Initialize a semaphore:

```
int sem_init(sem_t *sem, int pshared, unsigned value)
```

 sem: pointer to a sem_t structure allocated by user program


 pshared=0 if the sem is shared by threads in the same proc; otherwise, the sem can be shared by multiple processes (the sem_t structure should be allocated in the memory block shared by these processes)

 value: initial value of the sem


Destroy a semaphore:

```
int sem_destory(sem_t *sem)
```

Unnamed Semaphores

 Wait on a semaphore:

```
int sem_wait(sem_t *sem)
```

 Signal a semaphore:

```
int sem_post(sem_t *sem)
```


Named Semaphores

- ❏ Create and open a semaphore:

`sem_t *sem_open(char *name, O_CREAT, mode_t mode, unsigned value)`

- ❏ name: a permanent name (like a file name) for the sem

- ❏ mode: permission mode: combination of S_IRUSR (owner read), S_IWUSR (owner write), S_IXUSR (owner exec), S_IRGRP (owner's group read), S_IWGRP (group write), S_IXGRP (group exec), S_IROTH (other users read), S_IWOTH (other write), S_IXOTH (other exec)

- ❏ value: initial value of the sem

- ❏ Return the pointer to the sem_t structure of the sem

- ❏ Open an existing semaphore:

`sem_t *sem_open(char *name, O_EXCL)`

Named Semaphores

☞ Once a named semaphore has been open, it can be accessed using `sem_wait()` and `sem_post()` as an unnamed semaphore

☞ Close a semaphore:

```
sem_close(sem_t *sem);
```

☞ Destroy (permanently delete) a semaphore:

```
sem_unlink(char *name);
```

Classical Problems of Synchronization

-  Bounded-Buffer Problem
-  Readers and Writers Problem

Bounded-Buffer Problem

- ❏ N buffers, each can hold one item
- ❏ Two types of processes: producer processes and consumer processes
- ❏ Requirements:
 - ❏ No two processes can access the buffers simultaneously (mutual exclusion)
 - ❏ A producer cannot put an item to the buffers if the buffers are all full
 - ❏ A consumer cannot remove an item from the buffer if the buffers are all empty

Shared Variables:

```
int numAvailBuffer=N;
```

```
semaphore mutex; // =1, to ensure mutual exclusion
```

Producer:



```
wait(mutex);  
while(1) {  
    if(numAvailBuffer>0) {  
        numAvailBuffer --;  
        ... put data into a buffer ...  
        signal(mutex);  
        break;  
    } else {  
        signal(mutex);  
        wait(mutex);  
    }  
}
```

Consumer:



```
wait(mutex);
while(1){
    if(numAvailBuffer<N){
        numAvailBuffer++;
        ... fetch data from a buffer ...
        signal(mutex);
        break;
    }else{
        signal(mutex);
        wait(mutex);
    }
}
```

Bounded-Buffer Problem



To ensure mutual exclusion

-  Semaphore **mutex** initialized to the value 1 (number of admissible process)
-  Only when $\text{mutex} == 1$ can a process proceed to access the buffers (i.e., put in or remove an item)


To synchronize consumer processes

-  Semaphore **full** initialized to the value 0 (number of full buffers)
-  Only when $\text{full} > 0$ can a consumer process proceed

To synchronize producer processes

-  Semaphore **empty** initialized to the value N (number of empty buffers)
-  Only when $\text{empty} > 0$ can a producer process proceed

Bounded Buffer Problem

 The structure of the producer process

```
... produce an item  
wait (empty);  
wait (mutex);  
... add the item to the buffer  
signal (mutex);  
signal (full);
```

Bounded Buffer Problem

 The structure of the consumer process

```
wait (full);  
wait (mutex);  
... remove an item from  buffer  
signal (mutex);  
signal (empty);  
... consume the item in nextc
```

Reader-Writer Problem

- ❏ A data set is shared among a number of concurrent processes
 - ❏ Readers: only read the data set; they do **not** perform any updates
 - ❏ Writers: can both read and write
- ❏ Problem – allow multiple readers to read at the same time. Only one single writer can access the shared data at the same time.

Reader-Writer Problem: Analogies



It is fine for a group of people to watch (**read**) a TV simultaneously.



But changing (**write**) programs must be done sequentially (no concurrency)!

Reader-Writer Problem: Observations

- ❏ A writer is allowed to access the data set only when there is no other writer or reader accessing
- ❏ Readers can simultaneously access the data set
 - When a reader arrives, it is allowed to access the data set if (i) no one is accessing the data set, or (ii) some other reader is reading; otherwise (some writer is writing) , it should wait.
 - For a waiting reader, it is allowed to read if the writing writer completes.

Reader-Writer Problem: Observations → Solution

- ❏ A writer is allowed to access the data set only when there is no other writer or reader accessing
- ❏ Data Structure
 - ❏ Semaphore `wrt` initialized to 1 → to ensure at most one writer (no 2+ writers or one writer and reader(s)) allowed to access the data set (To guard access to data set)
- ❏ Code: The structure of a writer process

```
wait (wrt) ;  
.... writing is performed ...  
signal (wrt) ;
```

Reader-Writer Problem:


Observations → Solution

- ❏ Readers can simultaneously access the data set
 - When a reader arrives, it is allowed to access the data set if (i) no one is accessing the data set, or (ii) some other reader is reading; otherwise (some writer is writing) , it should wait.
 - For a waiting reader, it is allowed to read if the writing writer completes.

Additional Data Structures:


- ❏ Integer **readcount** initialized to 0 → to keep track of number of readers that are currently reading the data set
- ❏ Semaphore **mutex** initialized to 1 → to ensure mutual exclusive modification of readcount (To guard access to readcount)

Readers-Writers Problem: Solution

 The structure of a reader process

```
wait (mutex) ;  
if (readcount == 0) wait (wrt) ;  
    //check if there is already reader being reading  
    //if yes, go ahead to read;  
    //otherwise, go on only if no writer writing  
    //note: if this reader has to wait, follow-up readers should wait as well; so block them  
readcount ++ ;  
signal (mutex);  
    ... reading is performed...
```


Readers-Writers Problem: Solution

 The structure of a reader process

... reading is performed...

wait (mutex) ;


readcount - - ;


if (readcount == 0) signal (wrt) ; //if this is the last reader


signal (mutex) ;

Semaphore Implementation

 Must guarantee:

 No two processes can actively execute `wait()` on the same semaphore at the same time (but when one is blocked by wait, another is allowed to execute it)

 No two processes can modify the value of semaphore at the same time (i.e., `signal()` and “S- -” of `wait()` must be atomic)

 Thus, implementation becomes the critical section problem where the wait and signal code are placed in the critical section.

Semaphore Implementation

All competing processes share a variable lock initialized to false

```
Signal(S) {  
    while(testAndSet(&lock));  
    S++;  
    lock=False;  
}
```

Semaphore Implementation

All competing processes share a variable lock initialized to false

```
wait(S) {  
    while(testAndSet(&lock));  
    while (S <= 0) {  
        lock=False;  
        while(testAndSet(&lock);  
    }  
    S --;  
    lock=False;  
}
```

Semaphore Implementation

- ☐ Ideally, will not have **busy waiting** in critical section
- ☐ Otherwise, applications may spend lots of time in critical sections.

Semaphore Implementation with no Busy waiting


- ❏ With each semaphore, there are
 - ❏ value (of type integer)
 - ❏ list: an associated waiting queue. Each entry in a waiting queue has two items: (1) pointer to a process/thread; (2) pointer to next record in the list
- ❏ Two (atomic) operations:
 - ❏ **block** – place the process invoking the operation on a waiting queue.
 - ❏ **wakeup** – remove one of processes on a waiting queue and place it in the ready queue.

Semaphore Implementation with no Busy waiting

 (Rewritten) Definition of wait:

```
wait(semaphore *S) {  
    S->value--;  
    if (S->value < 0) {  
        add this process to S->list;  
        block();  
    }  
}
```

Semaphore Implementation with no Busy waiting

 Implementation of wait (example):


```
wait(semaphore *S) {  
    while(testAndSet(&lock));  
    S->value--;  
    if (S->value < 0) {  
        add this process to S->list;  
        lock=False;  
        block();  
    }  
    else lock=False;  
}
```


Semaphore Implementation with no Busy waiting

 Definition of signal:

```
signal(semaphore *S) {  
    S->value++;  
    if (S->value <= 0) {  
        remove a process P from S->list;  
        wakeup(P);  
    }  
}
```

Semaphore Implementation with no Busy waiting

 Implementation of signal (example):

```
signal(semaphore *S) {  
    while(testAndSet(&lock));  
    S->value++;  
    if (S->value <= 0) {  
        remove a process P from S->list;  
        wakeup(P);  
    }  
    lock=False;  
}
```